

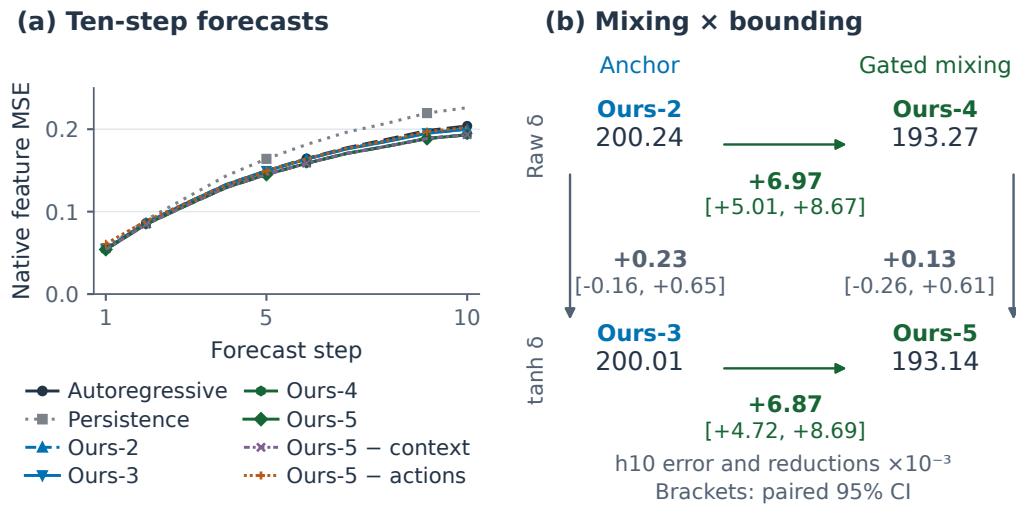


Figure 1: Candidate 1: 5.5 by 2.65 inches. Retains all current data marks and four node errors; compresses whitespace and uses a two-column, four-row legend.

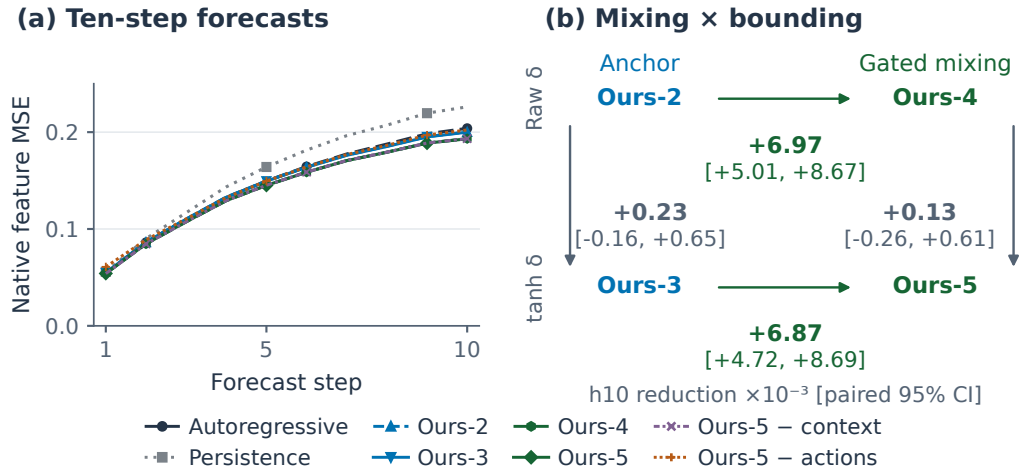


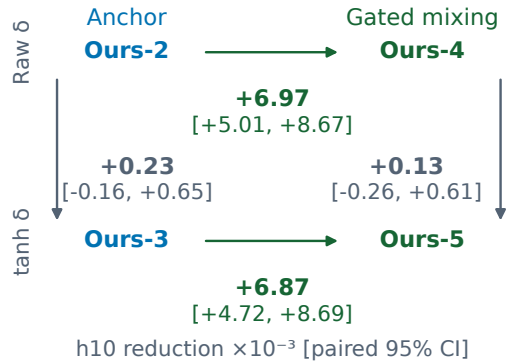
Figure 2: Candidate 2 (recommended): 5.5 by 2.60 inches. Native spatial prediction on original DROID validation. (a) Mean endpoint error at each of ten forecast steps; each consumes five native commands. (b) Arrows compare connected mechanisms at h10, showing error reductions and 95% session/seed bootstrap intervals. Mixing helps in both rows; incremental bounding and all interaction intervals include zero. Exact endpoint scores remain in Table 1; all 36 contrasts remain in the appendix. The follow-up uses revealed controls. Ours-1 uses a different protocol.

(a) Endpoint errors

Method	h5	h10
Autoregressive	0.149573	0.203950
Persistence	0.164028	0.226129
Ours-2	0.149606	0.200244
Ours-3	0.149347	0.200010
Ours-4	0.145037	0.193274
Ours-5	0.145074	0.193141
Ours-5 – context	0.145143	0.193330
Ours-5 – actions	0.149480	0.202249
Ours-5 vs AR	+3.01%	+5.30%

Native MSE; episode / seed means.

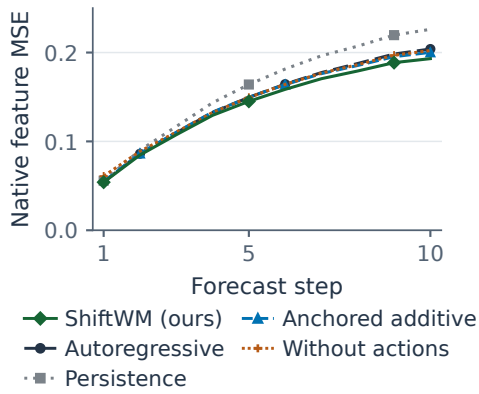
(b) Mixing × bounding



Trajectories would move to the appendix.

Figure 3: Candidate 3: 5.5 by 2.60 inches. Exact eight-row endpoint lookup plus component map; all forecasts would move to the appendix. This duplicates the separate main table, so it is not recommended when that table remains.

(a) Forecast error



(b) ShiftWM gains

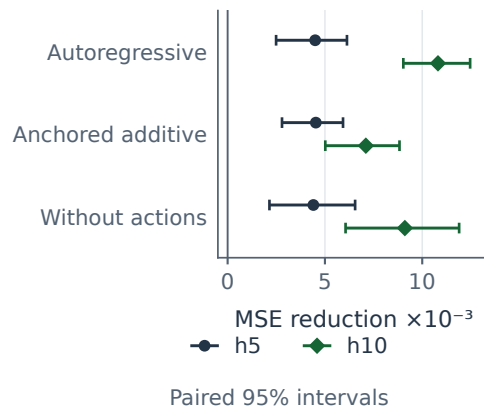
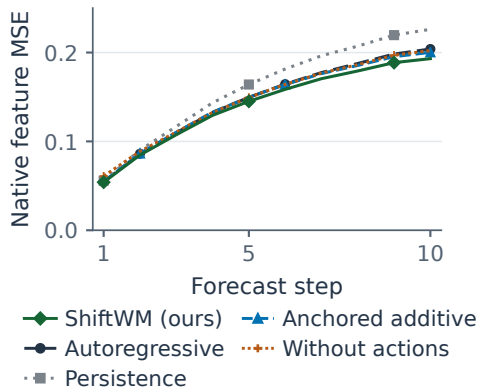


Figure 4: Candidate 4: 5.5 by 2.60 inches. ShiftWM (ours) on original DROID validation. (a) Mean endpoint feature error; each forecast step consumes five native commands. (b) Paired error reductions against autoregression, anchored additive prediction and the action-free ablation at steps 5 and 10. Error bars are 95% session/seed bootstrap intervals; connecting lines guide the eye. All 21 trained models completed 30 epochs. Other ablations remain in Table 1 and the appendix, including inconclusive context/bounding effects and unfavorable outcomes. The complete 36-contrast record remains unchanged; this is development evidence.

(a) Forecast error



(b) ShiftWM gains vs controls

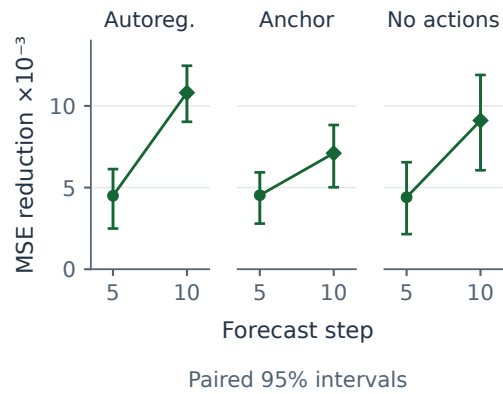


Figure 5: Candidate 5 (recommended main figure): 5.5 by 2.60 inches. ShiftWM (ours) on original DROID validation. (a) Mean endpoint feature error; each forecast step consumes five native commands. (b) Paired error reductions against autoregression, anchored additive prediction and the action-free ablation at steps 5 and 10. Error bars are 95% session/seed bootstrap intervals; connecting lines guide the eye. All 21 trained models completed 30 epochs. Other ablations remain in Table 1 and the appendix, including inconclusive context/bounding effects and unfavorable outcomes. The complete 36-contrast record remains unchanged; this is development evidence.