




Figure 1: Native spatial prediction on original DROID validation. (a) Mean endpoint error at each of ten forecast steps; each step consumes five native commands. (b) The factorial map separates gated mixing from innovation bounding at h10: nodes show error; arrows show paired reductions and 95% session/seed bootstrap intervals. Mixing helps in both rows; bounding intervals include zero. All interaction intervals also include zero. The follow-up reuses revealed controls; 21 models completed 30 epochs. Ours-1 uses a different protocol. Full metrics, all 36 contrasts, and scope details remain in the appendix.