


(a) Forecast error



(b) ShiftWM gains vs controls




Figure 1: ShiftWM (ours) on original DROID validation. (a) Mean endpoint feature error; each forecast step consumes five native commands. (b) Paired error reductions against autoregression, anchored additive prediction and the action-free ablation at steps 5 and 10. Error bars are 95% session/seed bootstrap intervals; connecting lines guide the eye. All 21 trained models completed 30 epochs. Other ablations remain in Table 1 and the appendix, including inconclusive context/bounding effects and unfavorable outcomes. The complete 36-contrast record remains unchanged; this is development evidence.