

(a) One held-out DROID window, two ways to forecast

frame t

forecast $t+10$ (decoded)

error at $t+10$

Direct: regenerates all 256 patches



every patch rewritten

static 0.16 moving 0.85

ShiftWM: moves the patches it has seen




learned transport

static 0.13 moving 0.72

feature error, shared scale: low  high  moving

(b) All 130 DROID test episodes



(c) Error reduction on every benchmark and as a plug-in

