


(a) Move, don't regenerate



DINOv2 patches Z_0
+ actions $\rightarrow Z_{1:10}$

Existing WMs regenerate every patch



ShiftWM moves what it has seen



$$\hat{Z}_k = (1-g)Z_0 \text{ keep} \\ + g\mathbf{T}_k \text{ move} \\ + \mathbf{r}_k \text{ correct}$$

✓ parallel from measured Z_0 ✓ plug-in head

(b) One held-out DROID window


true $t+10$ (decoded)



ShiftWM (ours)



Direct



AR (recursive)



AR-TF (DINO-WM-style)



Persistence (copy t)



Each tile: that model's own forecast, one shared decoder.
Corner: feature error at $t+10$ (lower is better).

(c) vs. best learned model

lower error than Direct in 100/100
horizon×motion bins (DROID)

